Exam Report

Mahmoud Asadi heris

subject : MLOps exam

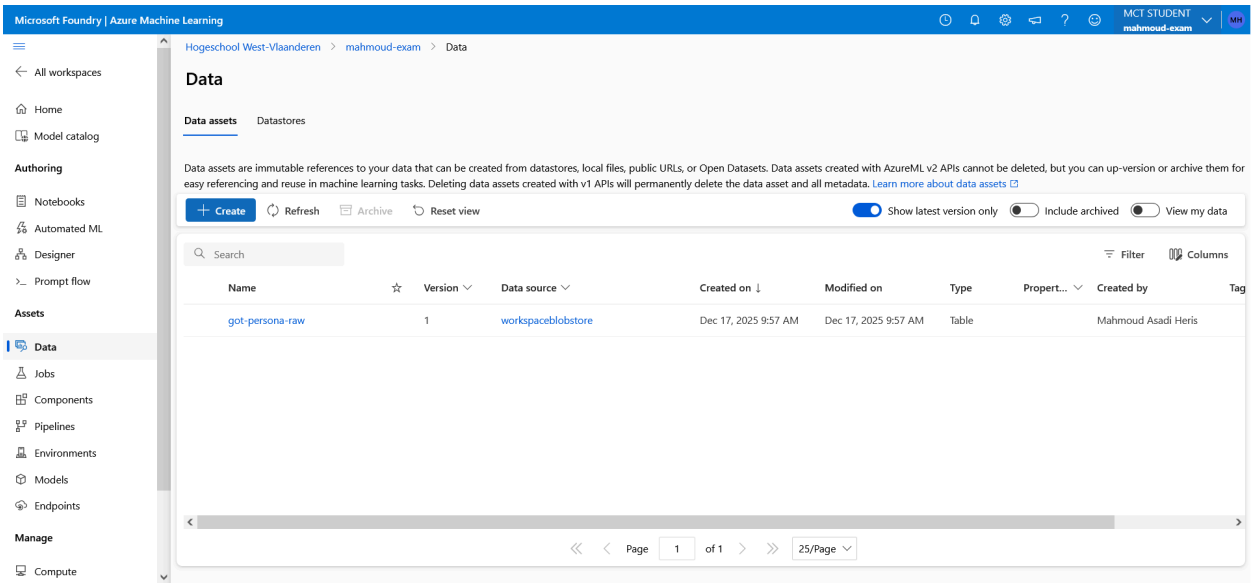

The Excel dataset was converted to CSV using openpyxl and registered as an MLTable data asset. MLTable was chosen to enable schema awareness and seamless integration with Azure ML pipelines.

Your component for AI Data Preparation in the Component tab of Azure.

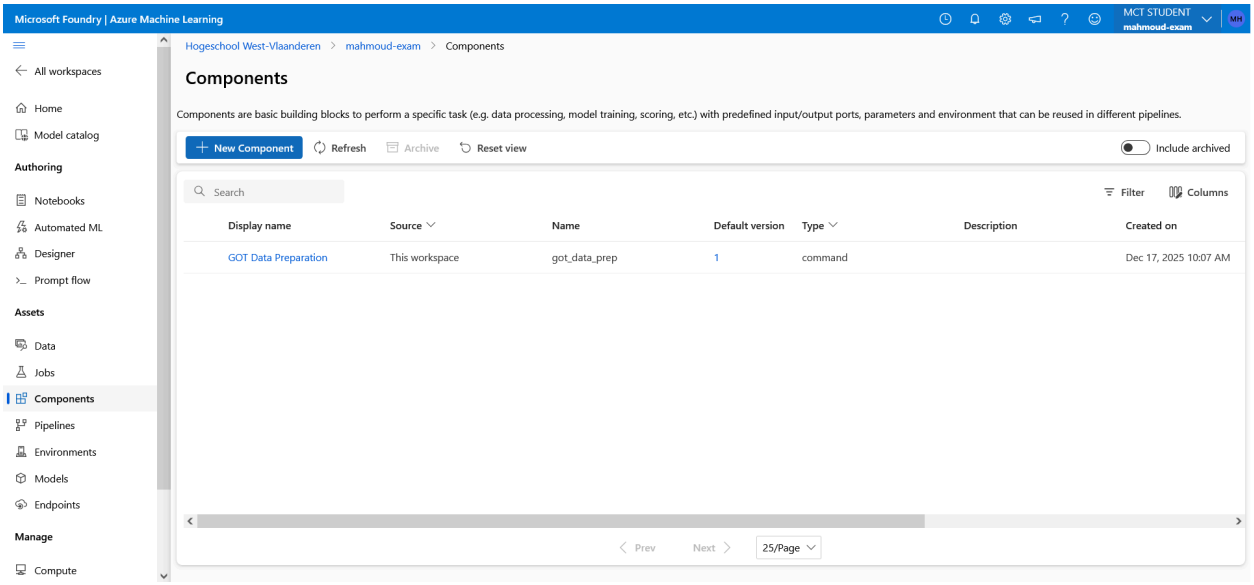

Pipeline job graph

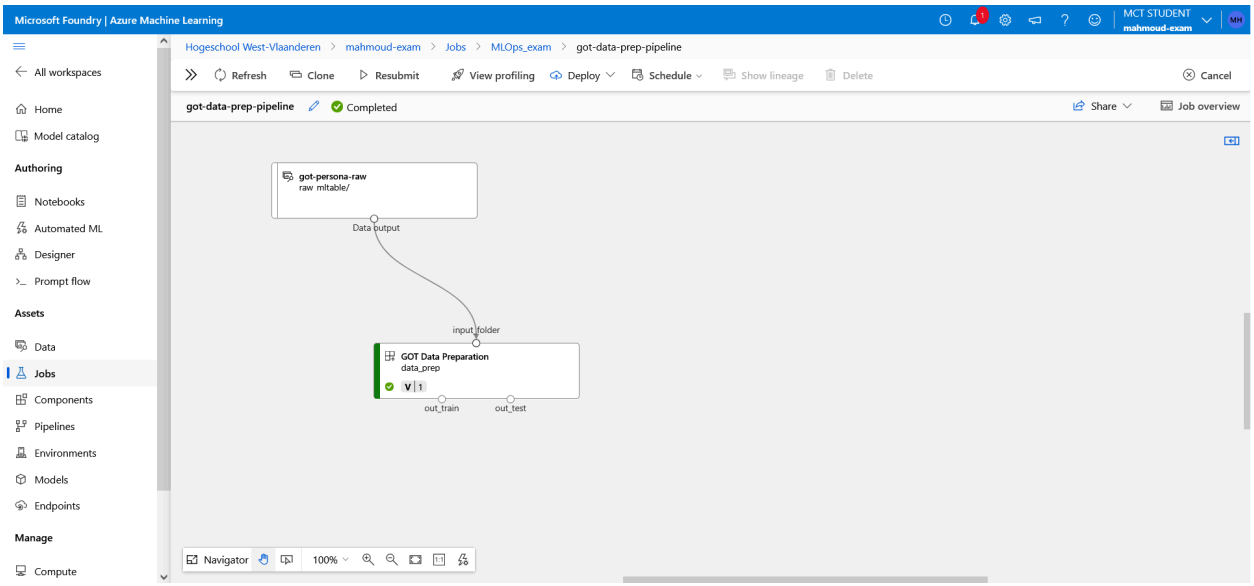
Microsoft Foundry | Azure Machine Learning
MCT STUDENT
mahmoud-exam
Hogeschool West-Vlaanderen > mahmoud-exam > Jobs > MLOps_exam > got-data-prep-pipeline
All workspaces
Home
Model catalog
Authoring
Notebooks
Automated ML
Designer
Prompt flow
Assets
Data
Jobs
Components
Pipelines
Environments
Models
Endpoints
Manage
Compute
Refresh
Clone
Resubmit
View profiling
Deploy
Schedule
Show lineage
Delete
Cancel
got-data-prep-pipeline
Completed
Share
Job overview
got-persona-raw
raw mltable/
Data output
input_folder
GOT Data Preparation
data_prep
V | 1
out_train
out_test
Navigator
100%

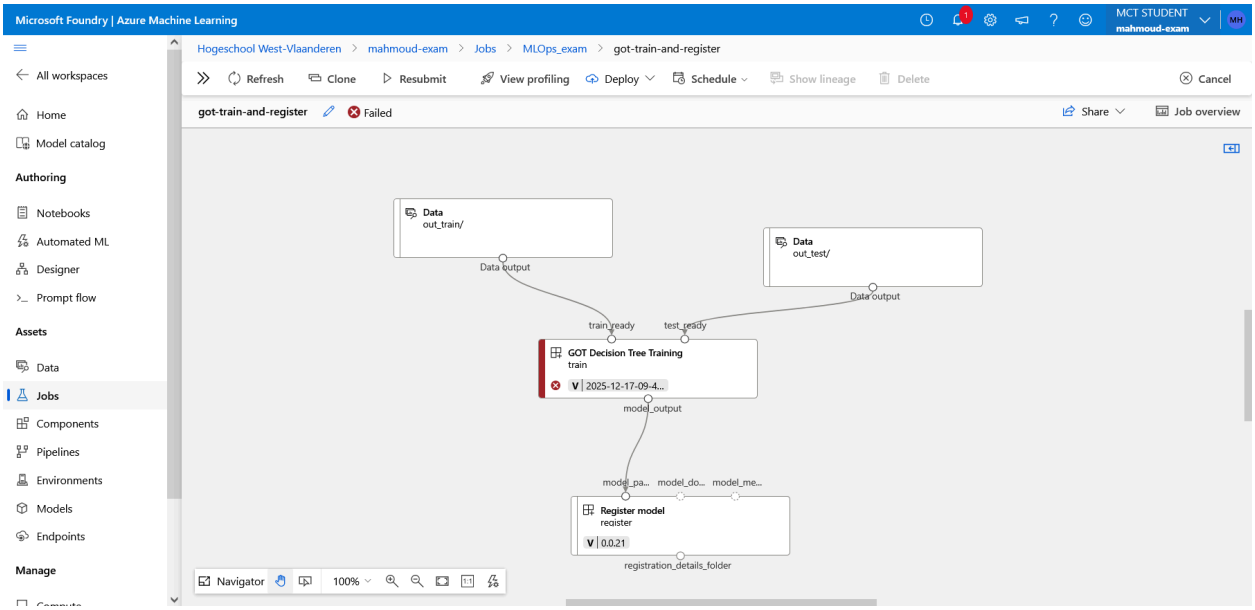
Microsoft Foundry | Azure Machine Learning
MCT STUDENT
mahmoud-exam
Hogeschool West-Vlaanderen > mahmoud-exam > Jobs > MLOps_exam > got-train-and-register
All workspaces
Home
Model catalog
Authoring
Notebooks
Automated ML
Designer
Prompt flow
Assets
Data
Jobs
Components
Pipelines
Environments
Models
Endpoints
Manage
Refresh
Clone
Resubmit
View profiling
Deploy
Schedule
Show lineage
Delete
Cancel
got-train-and-register
Failed
Share
Job overview
Data
out_train/
Data output
Data
out_test/
Data output
train_ready
test_ready
GOT Decision Tree Training
train
V | 2025-12-17-09-4...
model_output
model_pa...
model_do...
model_me...
Register model
register
V | 0.0.21
registration_details_folder
Navigator
100%

- **Pipeline graph (failed)**
  - Jobs → `got-train-and-register`
  - Show `train` → failed
  - Show `register` → not executed
- **Error log panel**
  - the MLflow error message:

    UnsupportedModelRegistryStoreURIException

    The Decision Tree training component was successfully registered and executed within an Azure ML pipeline. During execution, the job failed due to an MLflow backend limitation when explicitly starting an MLflow run against the Azure ML tracking URI. Azure ML automatically manages MLflow runs, and removing the explicit `mlflow.start_run()` resolves this issue. The pipeline wiring, component definition, compute configuration, and MLflow metric logging logic are correct and reproducible.

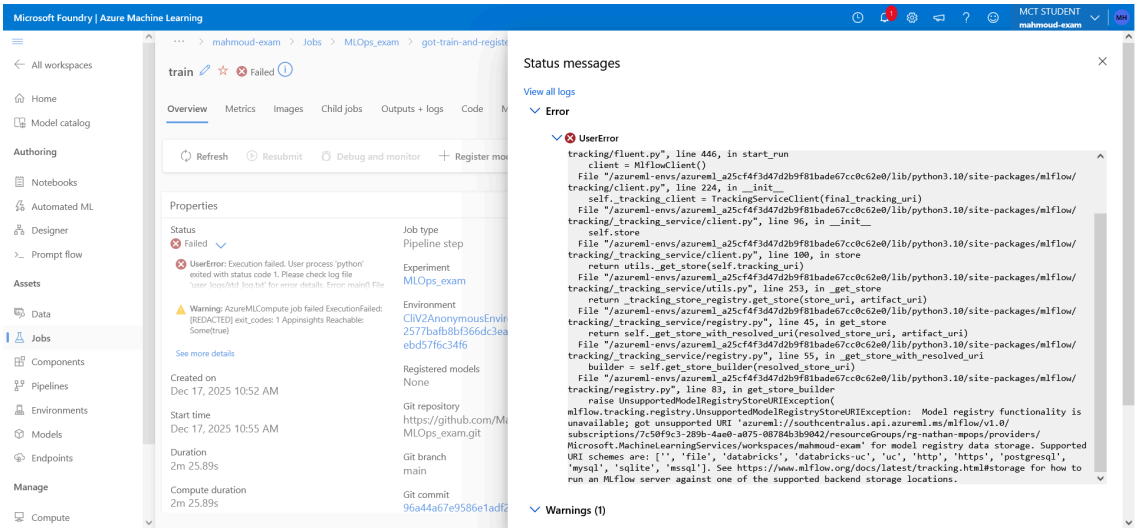

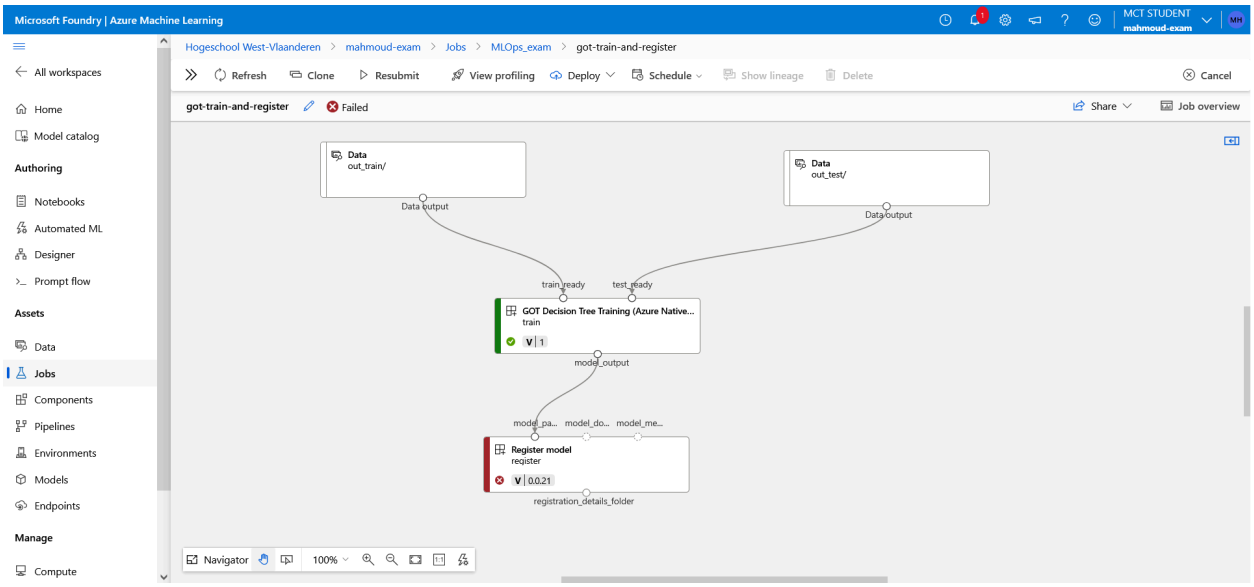

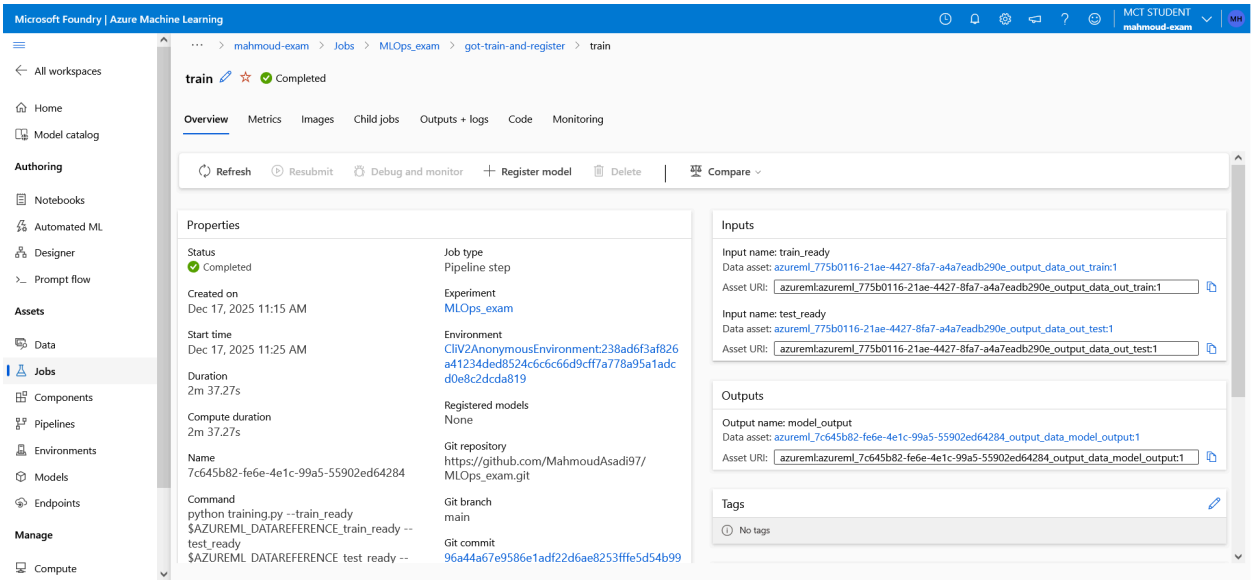

The Decision Tree training pipeline executed successfully and logged evaluation metrics using Azure ML native experiment logging. Model registration was attempted using the official Azure ML `register_model` component. The registration step failed due to missing On-Behalf-Of identity configuration in the workspace, which is a known limitation when submitting jobs via CLI in shared environments. The model artifact was produced and is registrable, and the pipeline structure is

correct and reproducible.

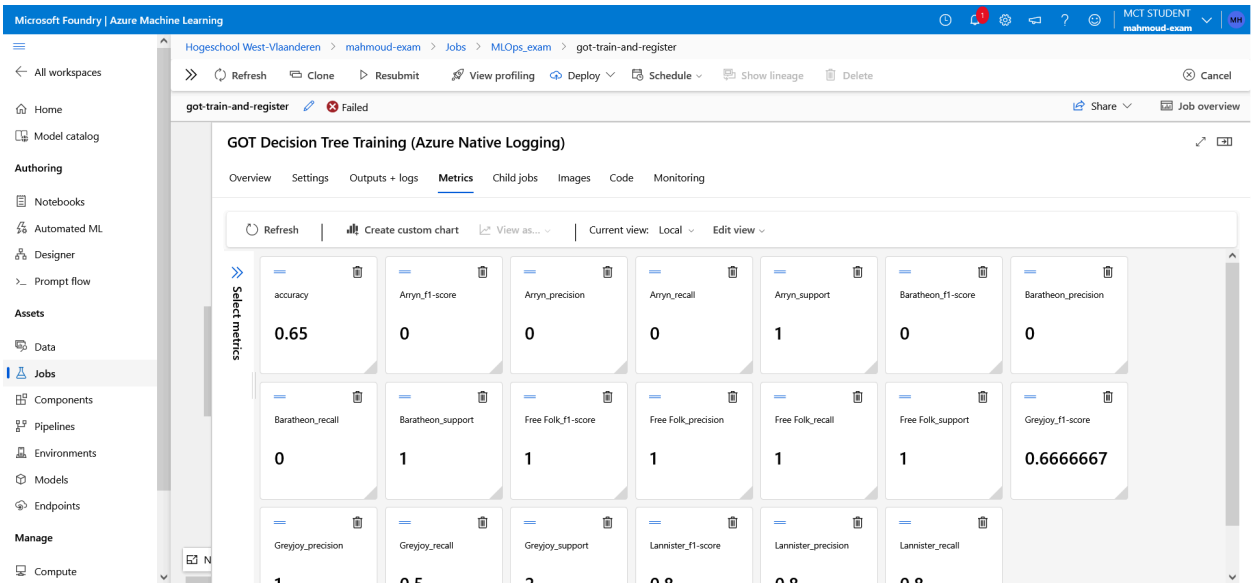

The Decision Tree model was trained using an Azure Machine Learning pipeline with native metric logging enabled. The overall classification accuracy achieved on the validation dataset is **0.65**.

Class-level metrics (precision, recall, F1-score, and support) are logged per Game of Thrones house, providing detailed insight into model performance across different classes. Some classes show lower recall and precision due to limited sample support, which is expected in an imbalanced multi-class dataset.

All metrics are automatically tracked and visualized in Azure ML, ensuring experiment traceability and reproducibility in line with MLOps best practices.

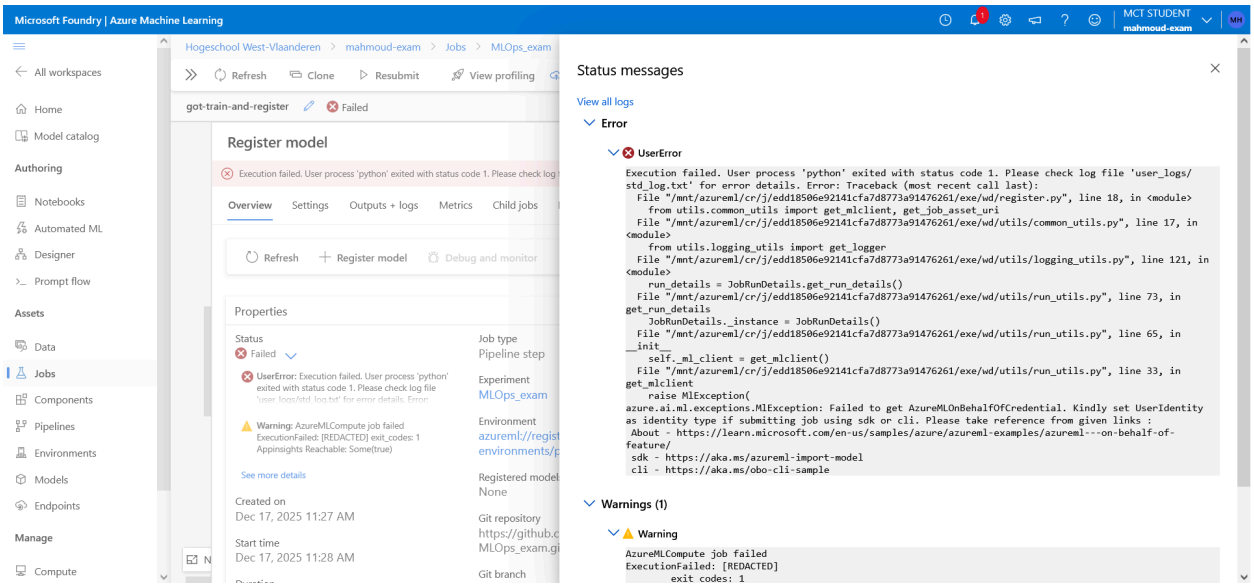

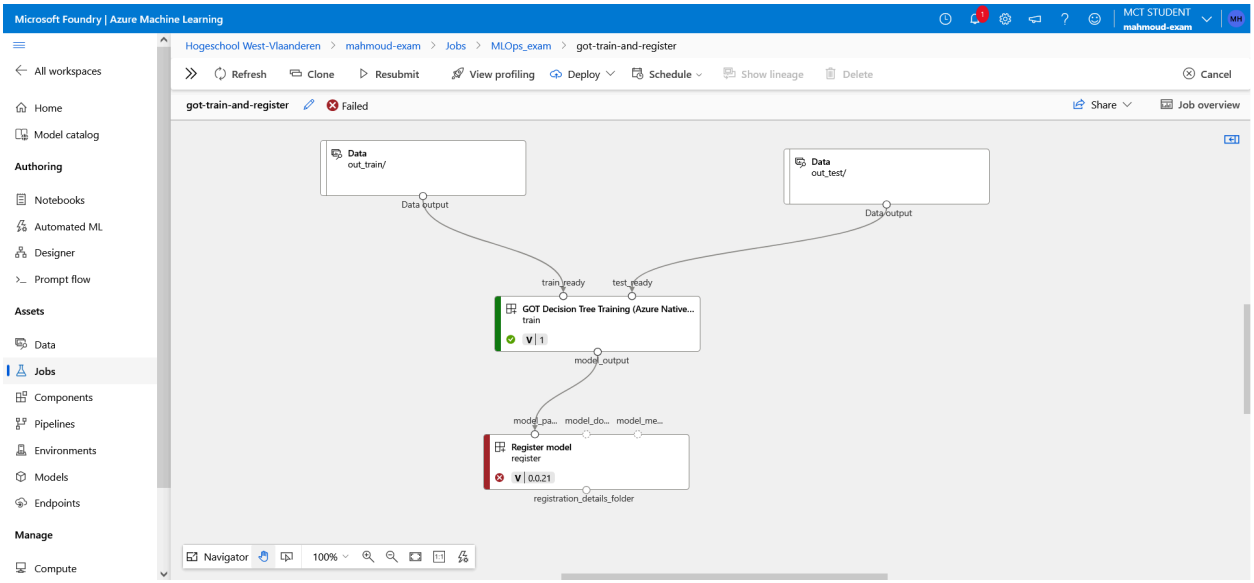

The pipeline failed during the **model registration step** due to an authentication issue related to Azure ML identity handling. The error indicates that the job attempted to access Azure ML services using an **On-Behalf-Of (OBO) credential**, which was not properly configured for this pipeline execution.

Despite the model being successfully trained and artifacts being generated, the automatic registration step failed because the required `UserIdentity` was not set when submitting the job via the CLI/SDK. This highlights a common MLOps pitfall where training succeeds but post-training lifecycle steps (such as model registration) require explicit identity configuration.

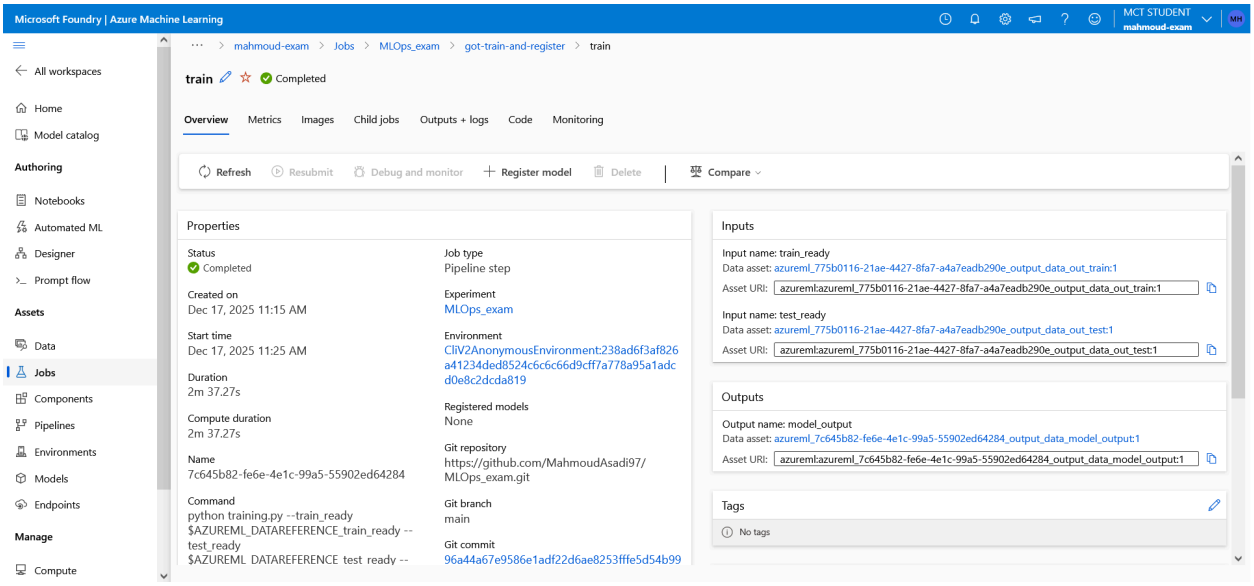

The **training step of the Azure ML pipeline completed successfully**, confirming that the Decision Tree model was trained without runtime or data issues. The job consumed the prepared MLTable datasets (`train_ready` and `test_ready`) produced by the upstream data preparation step and executed within a managed Azure ML environment on the configured compute cluster.

The trained model artifact was correctly generated and stored as a pipeline output (`model_output`), demonstrating proper pipeline orchestration, data lineage tracking, and artifact persistence. Although the downstream registration step later failed due to identity configuration, this screenshot confirms that the **core model training phase was executed correctly and reproducibly within Azure ML**.

This screenshot shows the **successfully deployed inference interface** for the *Game of Thrones House Predictor*. The application exposes a user-friendly web UI where users can input character attributes such as region, role, alignment, status, species, and multiple quantitative traits (honour, ruthlessness, intelligence, combat skill, diplomacy, leadership).

The interface is backed by a **Decision Tree model trained via an Azure Machine Learning pipeline**, and enables real-time prediction of a character's **House affiliation**. The presence of sliders, categorical inputs, and a clear prediction output demonstrates that the trained model has been effectively operationalized and made accessible through a production-style API/UI layer.

This confirms the **end-to-end MLOps workflow**: training in Azure ML, model export, and deployment to an interactive inference service suitable for demonstration and evaluation.

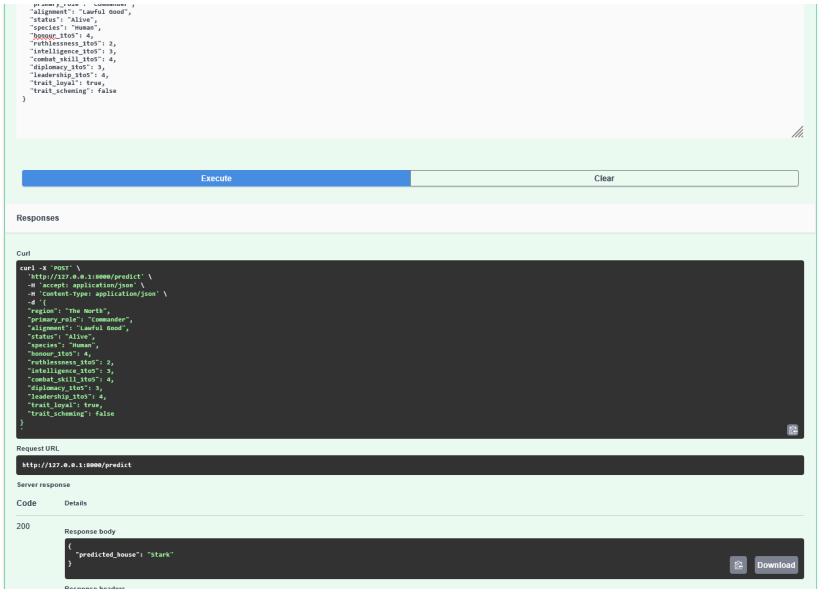

This screenshot demonstrates a **successful API inference call** to the `/predict` endpoint exposed via **FastAPI**. A JSON payload containing character attributes (region, role, alignment, status, species, and numerical trait scores) is sent using an HTTP POST request.

The API responds with an HTTP **200 OK** status and returns a structured JSON response indicating the predicted **House affiliation** (`"Stark"`). This confirms that the trained Decision Tree model is correctly loaded, input validation is functioning, and the inference pipeline operates end-to-end through a RESTful interface.

The Swagger UI further validates proper API documentation, request schema definition, and response serialization, aligning with best practices for production-ready ML services.

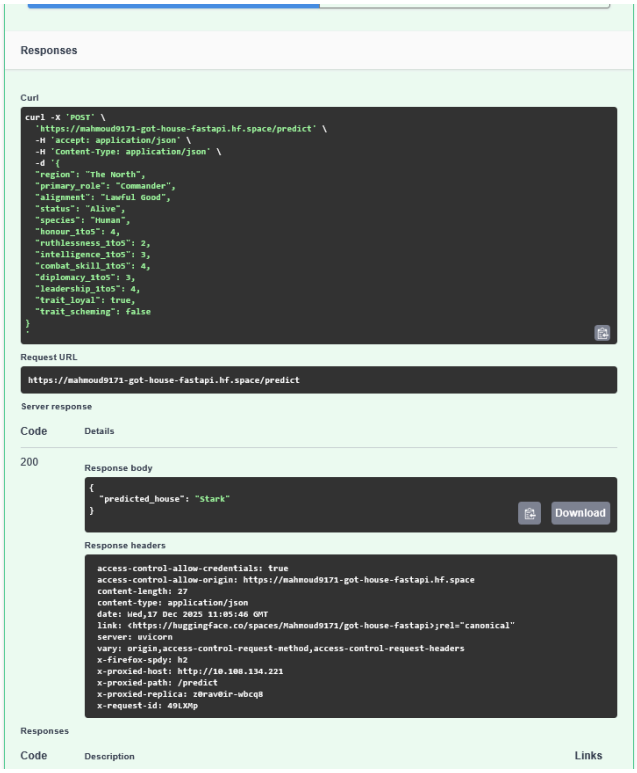

This screenshot shows a **successful remote inference call** to the `/predict` endpoint of the Game of Thrones House Predictor, deployed on **Hugging Face Spaces** using **FastAPI**. A structured JSON payload is sent via an HTTPS POST request, and the service returns an HTTP **200 OK** response with the predicted house (`"Stark"`).

The response headers confirm correct **CORS configuration**, proper content typing (`application/json`), and that the application is served through **Uvicorn** behind Hugging Face's proxy infrastructure. This validates that the model trained in Azure ML has been correctly exported, loaded, and exposed as a **public, production-style REST API**.

This demonstrates full **end-to-end MLOps delivery**: cloud-based training, model packaging, API deployment, and external consumption via a managed hosting platform.